

The Transformer model relies on self-attention.

Self-attention computes attention weights as $\text{softmax}(QK^T / \sqrt{d})$ V.

Figure 2 below shows an example attention heatmap.

